```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

▾ DESCRIPTIVE STATISTICS

```
import pandas as pd
df=pd.read_csv('/content/Churn Modelling.csv')
```

```
df.head()
```

➡

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	1	15634602	Hargrave	619	France	Female	42	2
1	2	15647311	Hill	608	Spain	Female	41	1
2	3	15619304	Onio	502	France	Female	42	8
3	4	15701354	Boni	699	France	Female	39	1
4	5	15737888	Mitchell	850	Spain	Female	43	2

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balar
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.0000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889:
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405:
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.0000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.0000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.5400
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.2400
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.0900

```
df.head(2)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenu:
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	

VISUALIZATIONS

```
sns.distplot(df.CreditScore)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: Futur
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fb607035b10>
```



```
sns.distplot(df.Age)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: Futur
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606f26dd0>
```



```
df.Gender.value_counts().plot(kind='barh')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb6069dc150>
```



```
df.Geography.value_counts().plot(kind='barh')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb60696b710>
```



```
df.Tenure.value_counts().plot(kind='barh')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606898550>
```



```
sns.distplot(df.EstimatedSalary)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: Futur
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fb6067b5d10>
```



```
sns.barplot(df.Geography, df.Exited)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWar
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606881750>
```



```
df.head(2)
```

|	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenu:
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	

```
sns.distplot(df.CreditScore)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: Futur
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fb60666b350>
```



```
sns.distplot(df.Age)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: Futur
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606569850>
```



```
df.Gender.value_counts().plot(kind='barh')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606469490>
```



```
df.Geography.value_counts().plot(kind='barh')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606f267d0>
```



```
df.Tenure.value_counts().plot(kind='barh')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb6067533d0>
```



```
sns.distplot(df.EstimatedSalary)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: Futur
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fb60638c450>
```

```
sns.barplot(df.Geography, df.Exited)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWar
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb6062c9d90>
```



```
df.head(2)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenu
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	

```
sns.barplot(x='Geography',y='Exited',data=df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb6062302d0>
```



```
sns.barplot(x='Gender',y='Exited',data=df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606219b90>
```



```
sns.barplot(x='NumOfProducts',y='Exited',data=df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb60620e650>
```



```
sns.barplot(x='HasCrCard',y='Exited',data=df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606189250>
```



```
sns.barplot(x='IsActiveMember',y='Exited',data=df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb6060cb6d0>
```



```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x7fb6060441d0>
```
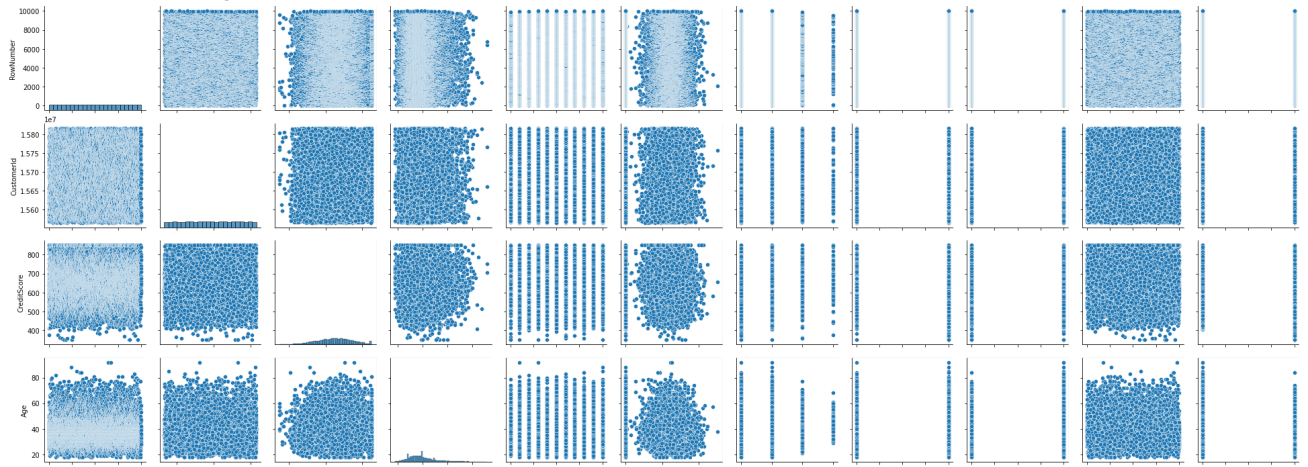


```
plt.figure(figsize=(8,5))
sns.heatmap(df.corr(),annot=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fb60071f650>
```



```
df.Exited.value_counts()
```

```
0    7963
1    2037
Name: Exited, dtype: int64
```

```
df.isnull().sum()
```

```
RowNumber            0
CustomerId           0
Surname              0
CreditScore          0
Geography            0
Gender               0
Age                  0
Tenure               0
```

```
Balance              0
NumOfProducts        0
HasCrCard            0
IsActiveMember       0
EstimatedSalary      0
Exited               0
dtype: int64
```

```
df.head(2)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	1	15634602	Hargrave	619	France	Female	42	2
1	2	15647311	Hill	608	Spain	Female	41	1

```
sns.boxplot(df.CreditScore)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb5febed650>
```



```
sns.boxplot(df.Age)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb5feba8890>
```



```
sns.boxplot(df.NumOfProducts)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb5feb24d90>
```



```
sns.boxplot(df.HasCrCard)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb606f1b5d0>
```



```
sns.boxplot(df.IsActiveMember)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb5febbab10>
```



```
sns.boxplot(df.EstimatedSalary)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb5fe9ec890>
```



```
sns.boxplot(df.Tenure)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
```

```
sns.boxplot(df.Balance)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb5fe944b50>
```


OUTLIERS

```python
def outlier_credit_score(df):
    IQR = df['CreditScore'].quantile(0.75) - df['CreditScore'].quantile(0.25)

    lower_range = df['CreditScore'].quantile(0.25) - (1.5 * IQR)
    upper_range = df['CreditScore'].quantile(0.75) + (1.5 * IQR)

    df.loc[df['CreditScore'] <= lower_range, 'CreditScore'] = lower_range
    df.loc[df['CreditScore'] >= upper_range, 'CreditScore'] = upper_range

outlier_credit_score(df)
```

```python
sns.boxplot(df.CreditScore)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb60088c990>
```

```python
def outlier_NOP(df):
    IQR = df['NumOfProducts'].quantile(0.75) - df['NumOfProducts'].quantile(0.25)

    lower_range = df['NumOfProducts'].quantile(0.25) - (1.5 * IQR)
    upper_range = df['NumOfProducts'].quantile(0.75) + (1.5 * IQR)

    df.loc[df['NumOfProducts'] <= lower_range, 'NumOfProducts'] = lower_range
    df.loc[df['NumOfProducts'] >= upper_range, 'NumOfProducts'] = upper_range

outlier_NOP(df)
```

```
        400         500         600         700         800
```

```python
sns.boxplot(df.NumOfProducts)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarnin
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb60090da10>
```



```python
def outlier_age(df):
    IQR = df['Age'].quantile(0.75) - df['Age'].quantile(0.25)

    lower_range = df['Age'].quantile(0.25) - (1.5 * IQR)
    upper_range = df['Age'].quantile(0.75) + (1.5 * IQR)

    df.loc[df['Age'] <= lower_range, 'Age'] = lower_range
    df.loc[df['Age'] >= upper_range, 'Age'] = upper_range

outlier_age(df)
```

```python
sns.boxplot(df.Age)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWar
  FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fb5fe888250>
```



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype
---  ------           --------------  -----
 0   RowNumber        10000 non-null  int64
 1   CustomerId       10000 non-null  int64
 2   Surname          10000 non-null  object
 3   CreditScore      10000 non-null  int64
 4   Geography        10000 non-null  object
 5   Gender           10000 non-null  object
 6   Age              10000 non-null  int64
 7   Tenure           10000 non-null  int64
 8   Balance          10000 non-null  float64
 9   NumOfProducts    10000 non-null  float64
 10  HasCrCard        10000 non-null  int64
 11  IsActiveMember   10000 non-null  int64
 12  EstimatedSalary  10000 non-null  float64
 13  Exited           10000 non-null  int64
dtypes: float64(3), int64(8), object(3)
memory usage: 1.1+ MB
```

```
df.head(2)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	1	15634602	Hargrave	619	France	Female	42	2
1	2	15647311	Hill	608	Spain	Female	41	1

```
df.drop(['CustomerId','RowNumber','Surname'],axis=1,inplace=True)
```

```
     --------------------------------------------------------------------
     KeyError                                    Traceback (most recent call last)
     <ipython-input-57-53894e3adb0e> in <module>
     ----> 1 df.drop(['CustomerId','RowNumber','Surname'],axis=1,inplace=True)
```

                            ⌄ 4 frames ─────
```
     /usr/local/lib/python3.7/dist-packages/pandas/core/indexes/base.py in
     drop(self, labels, errors)
        6015            if mask.any():
        6016                if errors != "ignore":
     -> 6017                    raise KeyError(f"{labels[mask]} not found in axis")
        6018                indexer = indexer[~mask]
        6019        return self.delete(indexer)
```

```
df.head(2)
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasC:
0	619	France	Female	42	2	0.00	1.0	
1	608	Spain	Female	41	1	83807.86	1.0	

```
from sklearn.preprocessing import LabelEncoder
le_geo = LabelEncoder()
le_gen = LabelEncoder()
df['Sex']=le_gen.fit_transform(df.Gender)
df['Country']=le_geo.fit_transform(df.Geography)
df.drop(['Geography','Gender'],axis=1,inplace=True)
```

```
df.head(2)
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	619	42	2	0.00	1.0	1	1
1	608	41	1	83807.86	1.0	0	1

```
X=df.drop('Exited',axis=1)
y=df.Exited
```

```
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X = sc.fit_transform(X)
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.2,
                                        random_state=42)
```

```
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
((8000, 10), (2000, 10), (8000,), (2000,))
```

Colab paid products - Cancel contracts here